Example Exploratory Data Analysis

This is an unmarked optional tutorial to show the kind of thinking that goes into an exploratory data analysis

The goal of this tutorial document is to walk through some of the common issues encountered in the early stages of an exploratory analysis on a set of data. It gives examples of common problem areas in:

- reading in data
- dealing with blanks
- dealing with factors

This data is a modified version of data from the New Zealand Election Survey, deliberately modified to introduce problems that occur naturally in many data sets.

Step One. Learn something about the data set.

In this case, the New Zealand Election Survey takes place every three years as a postal survey of a sample of registered electors. Some sampled electors were part of a sample panel of people surveyed at the previous election as part of a longitudinal study, others were randomly chosen from the electoral roll. Those electors that were part of the longitudinal panel group were randomly selected in previous elections.

As well as survey results, the data set includes information from the electoral roll, and weighting values for adjusting results. The full NZES data set has been reduced to a selected group of variables, making 3101 observations of 107 variables.

Step Two. Contemplate some questions.

Examining the codebook (or in this case the appendix at the end of the document to check out the variables).

For example, we might decide that since New Zealand is a Mixed Member Proportional voting system, where people get to vote for both an electorate (local) representative and a nationwide political party, that it would be interesting to look at strategic voting under conditions where there are many political parties to choose from. We identify some relevant variables of interest in the data, and investigate the nature of the individual variables before we explore their interactions. The kind of variables they are is going to shape our question.

Read in the data

There are many different kinds of data files in the world. Each one has its own issues when being read in by R. In this case the data is saved as a .RData file, which can be read in by using the `load()` command.

As with most reading in file commands, inside the parentheses needs to go a piece of text, in quotes, that is the path to the file from the working directory (the working directory is the folder that R is currently paying attention to). The easiest way to get a R Markdown (Rmd) Document and console cooperating about this is to place the file with the data in it in the same folder as the R Markdown (Rmd) Document, open the R Markdown (Rmd) document in RStudio so we are looking at the contents of the document in the editing window, then in the RStudio Session menu, use the **Set Working Directory - To Source File Location** command to make a common starting point. Then the code in the R Markdown (Rmd) document will use

the same working space regardless of whether we knit the document or run code chunks in the Console. In this case, if the `nzes2011.RData` file is in the same folder as the R Markdown (Rmd) document, hence it can be read into R with the following command:

```
load("selected_nzes2011.Rdata")
```

We also want to load packages that have functions in them we want to use. For this particular analysis we will only need the `dplyr` package, but for your project you will also likely need other packages as well, e.g. `ggplot2`.

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.0.5
```

Step Three. Prepare for the first question

As a first question, we might be interested in exploring the relationship between the party the person voted for, the party that was their favourite, and if they believed that their vote makes a difference – focusing on the question that are people who believe their vote makes a difference more likely to strategically vote for a party not their favourite. To achieve this, we familiarise ourselves with the variables `jpartyvote`, `jdiffvoting`, and `_singlefav`. First we check the codebook (see Appendix), then we explore the data.

Viewing the entire dataset in the Data Viewer window by clicking on the data frame's name in the Environment or running the `View()` command in the Console can be ineffective since the Data Viewer only shows the first 100 columns of the data frame.

Using the `str()` command on the entire dataset can also be equally ineffective. However we can subset the columns of interest and take a closer look at them. We can use the `dplyr` chain to select the variables of interest and investigate only their structure by adding `str()` at the end of the chain:

```
selected_nzes2011 %>%
  select(jpartyvote, jdiffvoting, _singlefav) %>%
  str()
```

If we try to run that line, we will get an error message about unexpected input or missing object.

We next need to diagnose where the problem lies – in the R code or in the data? The best way to troubleshoot this issue is to run each line of the `dplyr` chain one by one.

```
selected_nzes2011
```

The first line runs without any erros, but the second line gives an error

```
selected_nzes2011 %>%
  select(jpartyvote, jdiffvoting, _singlefav)
```

We know that `select()` is a valid `dplyr` function, so that cannot be the problem. This means the problem might be the variable names. The issue is that R has rules about what variable names are legal (e.g. no spaces, starting with a letter) and when data is loaded, R will often fix variable names to make them legal. This happened to the `_singlefav` at the time of loading the data.

We could check this by looking through every single variable name in the data with the `names()` command.

```
names(selected_nzes2011)
```

```
##   [1] "Jelect"        "jblogel"       "jnewspaper"    "jnatradio"
##   [5] "jtalkback"     "jdiscussp"     "jrallies"      "jpersuade"
##   [9] "jpcmoney"      "jpcposter"     "jlablike"      "jnatlike"
##  [13] "jgrnlike"      "jnzflike"      "jactlike"      "junflike"
##  [17] "jmaolike"      "jmnplike"      "jmostlike"     "jmostlikex"
##  [21] "jrepublic"     "jsphealth"     "jspedu"        "jspunemp"
##  [25] "jspdefence"    "jspsuper"      "jspbusind"     "jsppolice"
##  [29] "jspwelfare"    "jspenviro"     "jgovpdk"       "jgovplab"
##  [33] "jgovpnat"      "jgovpgrn"      "jgovpnzf"      "jgovpact"
##  [37] "jgovunf"       "jgovpmao"      "jgovpmnp"      "jnevervoteno"
##  [41] "jnevervotelab" "jnevervotenat" "jnevervotegrn" "jnevervotenzf"
##  [45] "jnevervoteact" "jnevervoteunf" "jnevervotemao" "jnevervotemnp"
##  [49] "jnevervoteoth" "jnevervoteothx" "jfirstpx"     "jsecondp"
##  [53] "jage"          "jlanguage"     "jlanguagex"    "jrollsex"
##  [57] "jhqual"        "jwkft"         "jwkpt"         "jwkun"
##  [61] "jwkret"        "jwkdis"        "jwksch"        "jwkunpo"
##  [65] "jwkunpi"       "jhhincome"     "jhhadults"     "jhhchn"
##  [69] "jmarital"      "r_jind"        "jlablr"        "jnatlr"
##  [73] "jgrnlr"        "jnzflr"        "jactlr"        "junflr"
##  [77] "jmaolr"        "jmnplr"        "jslflr"        "jrelservices"
##  [81] "jrelnone"      "jrelang"       "jrelpres"      "jrelcath"
##  [85] "jrelmeth"      "jrelbap"       "jrellat"       "jrelrat"
##  [89] "jrelfun"       "jrelothc"      "jrelnonc"      "jreligionx"
##  [93] "jreligiousity" "jethnicity_e"  "jethnicity_m"  "jethnicity_p"
##  [97] "jethnicity_a"  "jethnicity_o"  "jethnicityx"   "jethnicmost"
## [101] "jethnicmostx"  "jpartyvote"    "jelecvote"     "njptyvote"
## [105] "njelecvote"    "jdiffvoting"   "X_singlefav"
```

However, when we have hundreds of column names, a useful tip is to just search out only possible names. We can search the names for a fragment of the name by using the `grep("FRAGMENT", variable, value = TRUE)` command, which in this case might be:

```
grep("singlefav", names(selected_nzes2011), value = TRUE)
```

```
## [1] "X_singlefav"
```

The `value = TRUE` argument, as described in the help for the `grep()` function reports the mathing character string, as opposed to the index number for that string.

We can now confirm that the variable is called `X_singlefav`, so that is how we should be referring to it.

```
selected_nzes2011 %>%
  select(jpartyvote, jdiffvoting, X_singlefav) %>%
  str()
```

These are all categorical data, however they are recorded as characters (text strings) as opposed to factors.

An easy way of tabulating these data to see how many times each level of is to use the `group_by()` function along with the `summarise()` command:

```
selected_nzes2011 %>%
  group_by(jpartyvote) %>%
  summarise(count = n())
```

```
## # A tibble: 14 x 2
##    jpartyvote     count
##    <chr>          <int>
##  1 Act               29
##  2 ALC               10
##  3 Alliance           2
##  4 Another party      8
##  5 Conservative      74
##  6 Don't know        23
##  7 Green            348
##  8 Labour           749
##  9 Mana              62
## 10 Maori Party      128
## 11 National        1130
## 12 NZ First         216
## 13 United Future     14
## 14 <NA>             308
```

We can see that 23 people answered `"Don't know"`. Since our question is about people who knew which party they voted for, we might want to exclude these observations from our analysis. We can do so by `filter`ing them out.

```
selected_nzes2011 %>%
  filter(jpartyvote != "Don't know") %>%
  group_by(jpartyvote) %>%
  summarise(count = n())
```

```
## # A tibble: 12 x 2
##    jpartyvote     count
##    <chr>          <int>
##  1 Act               29
##  2 ALC               10
##  3 Alliance           2
##  4 Another party      8
##  5 Conservative      74
##  6 Green            348
##  7 Labour           749
##  8 Mana              62
##  9 Maori Party      128
## 10 National        1130
## 11 NZ First         216
## 12 United Future     14
```

Because there is a `%>%` at the end of the line, R knows to continue on to the next line, as with any other 'to be continued' symbol at the end of the line.

Note that adding the filter also got rid of the `NA` entries. NA (Not Available) is used to indicate blank entries – those observations for which there is no data recorded. It is always a good plan to be aware of NAs and deliberately include them in or exclude them from the analysis so that the final results are not surprising.

In this case since NA indicates that these people did not answer the question about which party they voted for, exluding them from the analysis makes sense.

We can also similarly view the levels and number of occurances of these levels in the `X_singlefav` variable:

```
selected_nzes2011 %>%
  group_by(X_singlefav) %>%
  summarise(count = n())
```

```
## # A tibble: 8 x 2
##   X_singlefav    count
##   <chr>          <int>
## 1 Act               33
## 2 Green            388
## 3 Labour          1043
## 4 Mana              47
## 5 National        1266
## 6 NZ First         138
## 7 United Future    128
## 8 <NA>              58
```

This set also has `NA` entries, but in this case we don't want to get rid of anything but the `NA`s so we need to target them directly. `NA` entries need special targeting because they do not actually exist (they are different to the text `"NA"` or a variable saved with the name `NA`).

If we only wanted to find the `NA`s we would use the `is.na()` function with the name of the variable inside the parentheses.

However since we want the entries that are **not** `NA`s we can use the **Not** operator, `!`, to indicate "we want all the ones that are not NA":`!is.na()`. Hence we can `filter` out all non NAs in our `dplyr` chain:

```
selected_nzes2011 %>%
  filter(!is.na(X_singlefav)) %>%
  group_by(X_singlefav) %>%
  summarise(count = n())
```

```
## # A tibble: 7 x 2
##   X_singlefav    count
##   <chr>          <int>
## 1 Act               33
## 2 Green            388
## 3 Labour          1043
## 4 Mana              47
## 5 National        1266
## 6 NZ First         138
## 7 United Future    128
```

And remember that we can `filter` for multiple characteristics at once:

```
selected_nzes2011 %>%
  filter(!is.na(X_singlefav), jpartyvote != "Don't know") %>%
  group_by(X_singlefav) %>%
  summarise(count=n())
```

```
## # A tibble: 7 x 2
##   X_singlefav    count
##   <chr>          <int>
## 1 Act               29
## 2 Green            354
## 3 Labour           914
## 4 Mana              42
## 5 National        1172
## 6 NZ First         119
## 7 United Future    115
```

If we examine the categories in `jdiffvoting` we can see that this variable has levels such as both `"Don't know"` and `NA`.

```
selected_nzes2011 %>%
  group_by(jdiffvoting) %>%
  summarise(count = n())
```

```
## # A tibble: 7 x 2
##   jdiffvoting                                                 count
##   <chr>                                                       <int>
## 1 Don't know                                                     63
## 2 Voting can make a big difference to what happens             1605
## 3 Voting can make a reasonable amount of difference to what happens   841
## 4 Voting can make some difference to what happens               339
## 5 Voting won't make any difference to what happens              119
## 6 Voting won't make much difference to what happens             106
## 7 <NA>                                                           28
```

We need to decide how we want to handle these levels in our analysis.

Remember that our main question is about whether people vote for their favorite party or a diffent one. Hence an straighforwrd approach would be to first determine whether each observation in the data represents a person who voted for the party same as their favorite party or different. This requires creating a new variable with the `mutate()` function.

In creating this variable we want to evaluate if for a given observation the values in the `jpartyvote` and `X_singlefav` variables are the same, or different:

```
selected_nzes2011 <- selected_nzes2011 %>%
  mutate(sameparty = ifelse(jpartyvote == X_singlefav, "same", "different"))
```

This creates a new variable named `sameparty` that has the value `"same"` if `jpartyvote` is equal to `X_singlefav`, and `"different"` otherwise.

We can again check our work by exploring the groupings in a View:

```
selected_nzes2011 %>%
group_by(jpartyvote, X_singlefav, sameparty) %>%
  summarise(count = n())
```

```
## `summarise()` has grouped output by 'jpartyvote', 'X_singlefav'. You can override using the `.groups
```

```
## # A tibble: 82 x 4
## # Groups:   jpartyvote, X_singlefav [82]
##     jpartyvote X_singlefav    sameparty count
##     <chr>      <chr>          <chr>     <int>
##  1 Act         Act            same         12
##  2 Act         Green          different     1
##  3 Act         National       different    14
##  4 Act         United Future  different     1
##  5 Act         <NA>           <NA>          1
##  6 ALC         Green          different     1
##  7 ALC         Labour         different     4
##  8 ALC         National       different     2
##  9 ALC         United Future  different     3
## 10 Alliance    Labour         different     1
## # ... with 72 more rows
```

We can see that observations where `jpartyvote` equaled `X_singlefav`, the value `"same"` was recorded for the new variable `sameparty`, and the value `"different"` was recorded otherwise. If either `jpartyvote` or `X_singlefav` had an NA, R could not check for equality and hence NA was recorded for the `sameparty` variable as well.

To view and summarize the "same" entries we can use the following:

```
selected_nzes2011 %>%
  group_by(jpartyvote, X_singlefav, sameparty) %>%
  summarise(count = n()) %>%
  filter(sameparty == "same")
```

```
## `summarise()` has grouped output by 'jpartyvote', 'X_singlefav'. You can override using the `.groups
```

```
## # A tibble: 7 x 4
## # Groups:   jpartyvote, X_singlefav [7]
##   jpartyvote     X_singlefav    sameparty count
##   <chr>          <chr>          <chr>     <int>
## 1 Act            Act            same         12
## 2 Green          Green          same        237
## 3 Labour         Labour         same        632
## 4 Mana           Mana           same         31
## 5 National       National       same       1004
## 6 NZ First       NZ First       same         82
## 7 United Future  United Future  same          5
```

And to view and summarize the "different" entries we can use the following:

```
selected_nzes2011 %>%
  group_by(jpartyvote, X_singlefav, sameparty) %>%
  summarise(count = n()) %>%
  filter(sameparty == "different")
```

```
## `summarise()` has grouped output by 'jpartyvote', 'X_singlefav'. You can override using the `.groups
```

```
## # A tibble: 59 x 4
```

```
## # Groups:   jpartyvote, X_singlefav [59]
##    jpartyvote    X_singlefav   sameparty count
##    <chr>         <chr>         <chr>     <int>
##  1 Act           Green         different     1
##  2 Act           National      different    14
##  3 Act           United Future different     1
##  4 ALC           Green         different     1
##  5 ALC           Labour        different     4
##  6 ALC           National      different     2
##  7 ALC           United Future different     3
##  8 Alliance      Labour        different     1
##  9 Alliance      National      different     1
## 10 Another party Green         different     2
## # ... with 49 more rows
```

We can also check how we got any `NAs` we have by using the `is.na()` function:

```r
selected_nzes2011 %>%
  group_by(jpartyvote, X_singlefav, sameparty) %>%
  summarise(count = n()) %>%
  filter(is.na(sameparty))
```

```
## 'summarise()' has grouped output by 'jpartyvote', 'X_singlefav'. You can override using the '.groups
```

```
## # A tibble: 16 x 4
## # Groups:   jpartyvote, X_singlefav [16]
##    jpartyvote   X_singlefav   sameparty count
##    <chr>        <chr>         <chr>     <int>
##  1 Act          <NA>          <NA>          1
##  2 Conservative <NA>          <NA>          1
##  3 Don't know   <NA>          <NA>          7
##  4 Green        <NA>          <NA>          1
##  5 Labour       <NA>          <NA>         11
##  6 Maori Party  <NA>          <NA>          2
##  7 National     <NA>          <NA>          7
##  8 NZ First     <NA>          <NA>          2
##  9 <NA>         Act           <NA>          4
## 10 <NA>         Green         <NA>         32
## 11 <NA>         Labour        <NA>        121
## 12 <NA>         Mana          <NA>          4
## 13 <NA>         National      <NA>         92
## 14 <NA>         NZ First      <NA>         17
## 15 <NA>         United Future <NA>         12
## 16 <NA>         <NA>          <NA>         26
```

The checks show that the observations with `NAs` in the `sameparty`are going to be excluded from the analysis when we fiter out the `NAs` in the `jpartyvote` and `X_singlefav` variables, so we don't need to worry about them anymore.

Step four. Prepare for the second question

As a second question, we might be interested in exploring the relationship between age of voters and how much they like the NZ First party. We become familiar with the variables `jnzflike` and `jage` in the codebook, then explore the data.

```r
str(selected_nzes2011$jnzflike)
```

```
##  Factor w/ 12 levels "0","1","10","2",..: 1 1 4 10 4 11 NA NA 1 12 ...
```

```r
str(selected_nzes2011$jage)
```

```
##  int [1:3101] 37 37 28 71 43 NA 59 68 64 70 ...
```

`jnzflike` is a factor variable, in fact it's ordinal and by default the levels are listed in alphabetical order. Since this is a categorical variable, we can also summarize the occurances of each level with `group_by()` and `summarise()` again:

```r
selected_nzes2011 %>%
  group_by(jnzflike) %>%
  summarise(count = n())
```

```
## # A tibble: 13 x 2
##    jnzflike    count
##    <fct>       <int>
##  1 0             622
##  2 1             298
##  3 10            134
##  4 2             266
##  5 3             227
##  6 4             162
##  7 5             544
##  8 6             165
##  9 7             138
## 10 8             107
## 11 9              81
## 12 Don't know    224
## 13 <NA>          133
```

While `jnzflike` is on a 0 to 10 scale, this variable also has a level labeled `"Don't know"`, which is why R stores this variable as not a numeric variable.

`jage`, on the other hand, is an integer, with values that are whole numbers between 0 and infinity (or `NA`). For this variable we would want to take a look at numerical summaries such as means, medians, etc.

```r
selected_nzes2011 %>%
  summarise(agemean = mean(jage), agemedian = median(jage), agesd = sd(jage),
            agemin = min(jage), agemax = max(jage))
```

```
##   agemean agemedian agesd agemin agemax
## 1      NA        NA    NA     NA     NA
```

What went wrong? The reason why all of the results were reported as NAs is that there were some NA entries in the `jage` variable (people not reporting their age). Since it is not possible to take the average of a series of values that contain `NA`s, obtaining the numerical summaries requires that we exclude the `NA`s from the calculation.

Most numerical summary functions allow us to easily exclude `NA`s with the `na.rm` argument. See the help documentation for the `median` function for more information.

```
?median
```

```
## starting httpd help server ... done
```

An alternative approach is just to `filter` out the NAs first, and then ask for the numerical summaries:

```r
selected_nzes2011 %>%
  filter(!(is.na(jage))) %>%
  summarise(agemean = mean(jage), agemedian = median(jage), agesd = sd(jage),
            agemin = min(jage), agemax = max(jage))
```

```
##    agemean agemedian   agesd agemin agemax
## 1 53.22328        54 17.5371     18    100
```

An age range of 18 to 100 is a reasonable age range for a voting age population, so there are no obvious errors in the data. If there were, we would need to decide if we should filter them out of the analysis.

Having gained some familiarity with the specific variables we are using, we next need to consider if there is additional work we should do on the data in investigating the question. There are a number of different approaches we might take. For example, we could consider if those that strongly like NZ First are older than those that strongly dislike NZ First, or we could consider if old people like NZ First more than young people.

Approach 1: Strongly liking and disliking NZ First and age

If we wanted to select only two of the possible levels in how much people like NZ First, we can filter for these specific levels. When interested in filtering for multiple values a variable can take, the `%in%` operator can come in handy:

```r
selected_nzes2011 %>%
  filter(jnzflike %in% c("0","10")) %>%
  group_by(jnzflike) %>%
  summarise(count = n())
```

```
## # A tibble: 2 x 2
##   jnzflike count
##   <fct>    <int>
## 1 0          622
## 2 10         134
```

Remember that the `jnzflike` is not a numerical variable, hence we use the quotation marks around the values (even though they happen to be numbers).

This is an example of simpligying the analysis by considering only two levels of a categorical variable, as opposed to all possible levels.

Approach 2: Age and liking for NZ First

We might also like to refine our question slightly, asking do people above retirement age (65 in New Zealand) like NZ First more than younger people. To do this we can turn the numeric age variable into a categorical variable based on whether people are 65 years or older or younger than 65. Once again we make use of the `mutate()` and `ifelse()` functions:

```
selected_nzes2011 <- selected_nzes2011 %>%
  mutate(retiredage = ifelse(jage >= 65, "retired age", "working age"))
selected_nzes2011 %>%
  group_by(retiredage) %>%
  summarise(count = n())
```

```
## # A tibble: 3 x 2
##   retiredage  count
##   <chr>       <int>
## 1 retired age   876
## 2 working age  2156
## 3 <NA>           69
```

We can see that individuals in the dataset are now labeled as either `"retired age"` or `"working age"` or neither (`NA`), which we can easily filter out if need be.

This is an example of using a numerical threshold to convert a numerical variable to a categorical variable.

For approach 2, we might also be want to turn the scale of liking into numeric values, because at the moment we cannot easily get summary information of the data in factor form. For example, if we ty to run the following command, we get an error saying "need numeric data".

```
selected_nzes2011 %>%
  group_by(retiredage) %>%
  summarise(medlike = median(jnzflike))
```

it generates a "need numeric data" error.

We can change the type of data with functions of the form `as.thingtochangeto()`, but it is easy to go wrong with factors. For example, this is wrong:

```
selected_nzes2011 <- selected_nzes2011 %>%
  mutate(numlikenzf = as.numeric(jnzflike))
```

We can see it has gone wrong if we use grouping to check our work (and it is a very good plan to check our work after converting factors).

```
selected_nzes2011 %>%
  group_by(jnzflike, numlikenzf) %>%
  summarise(count = n())
```

```
## `summarise()` has grouped output by 'jnzflike'. You can override using the `.groups` argument.
```

```
## # A tibble: 13 x 3
## # Groups:   jnzflike [13]
##    jnzflike  numlikenzf count
##    <fct>          <dbl> <int>
## 1 0                  1   622
## 2 1                  2   298
## 3 10                 3   134
## 4 2                  4   266
## 5 3                  5   227
## 6 4                  6   162
```

```
##  7 5                         7   544
##  8 6                         8   165
##  9 7                         9   138
## 10 8                        10   107
## 11 9                        11    81
## 12 Don't know               12   224
## 13 <NA>                     NA   133
```

Factor entries have two parts: the text we see on the screen, and a numeric order (remember how 10 was coming between 1 and 2 because of the alphabetical order). When we say "turn this into a number", R uses the numeric order in which it stores the values to do that conversion, as opposed to the names of the levels of the categorical variable. Hence, we need a conversion method that will use the text strings that label the levels, as opposed to the storage order of these levels. We can do this by first saving the variable as a character variable, and then turning it into a number:

```
selected_nzes2011 <- selected_nzes2011 %>%
  mutate(numlikenzf = as.numeric(as.character(jnzflike)))
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

The warning "NAs introduced by coercion" happens since the level `"Don't know"` cannot be turned into a number. But this should be fine for our purposes since we are interested in the numerical responses anyway.

```
selected_nzes2011 %>%
  group_by(jnzflike, numlikenzf) %>%
  summarise(count = n())
```

```
## `summarise()` has grouped output by 'jnzflike'. You can override using the `.groups` argument.
```

```
## # A tibble: 13 x 3
## # Groups:   jnzflike [13]
##    jnzflike    numlikenzf count
##    <fct>            <dbl> <int>
##  1 0                    0   622
##  2 1                    1   298
##  3 10                  10   134
##  4 2                    2   266
##  5 3                    3   227
##  6 4                    4   162
##  7 5                    5   544
##  8 6                    6   165
##  9 7                    7   138
## 10 8                    8   107
## 11 9                    9    81
## 12 Don't know          NA   224
## 13 <NA>                NA   133
```

Converting the factor to a character first ensures that the numerical values used in the labels of the levels of the categorical variable are used.

Now that we cleaned up the data in a way that addresses the needs of the research questions we want to explore, we are ready to continue with our analysis.

Appendix: List of fields in example data

Variable	Question	DataType
jactlike	A14: how much like Act	Factor
jactlr	A18: Act on left-right scale	chr
jage	Respondent's age in years	int
jblogel	A6h: visit political blog for election	chr
jdiffvoting	A13: does voting make any difference to what happens	chr
jdiscussp	A11a: how often discussed politics with others	chr
Jelect	Electorate	int
jelecvote	C4: if cast electorate vote, for which party's candidate	chr
jethnicity_a	F19a: ethnicity - Asian	chr
jethnicity_e	F19a: ethnicity - NZ European	chr
jethnicity_m	F19a: ethnicity - NZ Maori	chr
jethnicity_o	F19a: ethnicity - Other	chr
jethnicity_p	F19a: ethnicity - Pacific	chr
jethnicityx	F19ax: other ethnic group belonged to detail	chr
jethnicmost	F19b: ethnic group identified with most	chr
jethnicmostx	F19bx: other ethnic group identified with most	chr
jfirstpx	C10x: on election day other party most wanted to be in government	chr
jgovpact	C8: Act helped form the government after 2008 election	chr
jgovpdk	C8: can't recall which parties formed the government after 2008 election	chr
jgovpgrn	C8: Greens helped form the government after 2008 election	chr
jgovplab	C8: Labour helped form the government after 2008 election	chr
jgovpmao	C8: Maori Party helped form the government after 2008 election	chr
jgovpmnp	C8: Mana Party helped form the government after 2008 election	chr
jgovpnat	C8: National helped form the government after 2008 election	chr
jgovpnzf	C8: NZ First helped form the government after 2008 election	chr
jgovunf	C8: United Future helped form the government after 2008 election	chr
jgrnlike	A14: how much like Greens	Factor
jgrnlr	A18: Greens on left-right scale	chr
jhhadults	F23a: number of adults in household	int
jhhchn	F23b: number of children in household	int
jhhincome	F22: household income between 1 April 2010 and 31 March 2011	chr
jhqual	F8: highest formal educational qualification	chr
jlablike	A14: how much like Labour	Factor
jlablr	A18: Labour on left-right scale	chr
jlanguage	F3: main language spoken at your home	chr
jlanguagex	F3x: other main language spoken	chr
jmaolike	A14: how much like Maori Party	Factor
jmaolr	A18: Maori Party on left-right scale	chr
jmarital	F24: marital status	chr
jmnplike	A14: how much like Mana Party	Factor
jmnplr	A18: Mana Party on left-right scale	chr
jmostlike	A15: on election day which party liked most	chr
jmostlikex	A15x: other party liked most	chr
jnatlike	A14: how much like National	Factor
jnatlr	A18: National on left-right scale	chr
jnatradio	A10d: how often followed election news on Radio New Zealand: National	chr
jnevervoteact	C16: would never vote for Act	chr
jnevervotegrn	C16: would never vote for Greens	chr
jnevervotelab	C16: would never vote for Labour	chr
jnevervotemao	C16: would never vote for Maori Party	chr
jnevervotemnp	C16: would never vote for Mana Party	chr
jnevervotenat	C16: would never vote for National	chr

Variable	Question	DataType
jnevervoteno	C16: no party for which you would never vote	chr
jnevervotenzf	C16: would never vote for NZ First	chr
jnevervoteoth	C16: would never vote for another party	chr
jnevervoteothx	C16: other party for for which you would never vote	chr
jnevervoteunf	C16: would never vote for United Future	chr
jnewspaper	A10c: how often followed election news in newspaper	chr
jnzflike	A14: how much like NZ First	Factor
jnzflr	A18: NZ First on left-right scale	chr
jpartyvote	C3: if cast party vote, for which party	chr
jpcmoney	A11d: how often contributed money to a party or candidate	chr
jpcposter	A11e: how often put up party or candidate posters	chr
jpersuade	A11c: how often talk to anyone to persuade them how to vote	chr
jrallies	A11b: how often attended political meetings or rallies	chr
jrelang	F17: anglican	chr
jrelbap	F17: baptist	chr
jrelcath	F17: catholic	chr
jrelfun	F17: independent-fundamentalist-pentecostal church	chr
jreligionx	F17x: other religion detail	chr
jreligiousity	F18: how religious are you	chr
jrellat	F17: latter day saints	chr
jrelmeth	F17: methodist	chr
jrelnonc	F17: non-Christian	chr
jrelnone	F17: no religion	chr
jrelothc	F17: other Christian	chr
jrelpres	F17: presbyterian	chr
jrelrat	F17: ratana	chr
jrelservices	F16: apart from weddings, funerals, baptisms, how often do you attend religious services	chr
jrepublic	B1: should NZ become a republic or retain Queen as head of state	chr
jrollsex	Respondent's gender from electoral roll	chr
jsecondp	C11: on election day which party overall was you second choice to be in government	chr
jslflr	A19: yourself on left-right scale	chr
jspbusind	B3f: should there be more or less public spending on business and industry	chr
jspdefence	B3d: should there be more or less public spending on defence	chr
jspedu	B3b: should there be more or less public spending on education	chr
jspenviro	B3i: should there be more or less public spending on the environment	chr
jsphealth	B3a: should there be more or less public spending on health	chr
jsppolice	B3g: should there be more or less public spending on police and law enforcement	chr
jspsuper	B3e: should there be more or less public spending on superannuation	chr
jspunemp	B3c: should there be more or less public spending on unemployment benefits	chr
jspwelfare	B3h: should there be more or less public spending on welfare benefits	chr
jtalkback	A10e: how often followed election news on talkback radio	chr
junflike	A14: how much like United Future	Factor
junflr	A18: United Future on left-right scale	chr
jwkdis	F9: disabled, unable to work	chr
jwkft	F9: working full-time for pay or other income	chr
jwkpt	F9: working part-time for pay or other income	chr
jwkret	F9: retired	chr
jwksch	F9: at school, university, or other educational institution	chr

Variable	Question	DataType
jwkun	F9: unemployed, laid off, looking for work	chr
jwkunpi	F9: working unpaid within the home	chr
jwkunpo	F9: working unpaid outside the home	chr
njelecvote	Electorate Vote with nonvote	chr
njptyvote	Party Vote with nonvote	chr
r_jind	Respondent Industry Codes	chr
_singlefav	Caluclated Variable of most liked of major parties Question A14	chr